



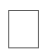


Best Practices for Reliability and Availability in Modern Data Management.

Abhishek Mishra, Developer Advocate, Yugabyte



\$ whoami: **Abhishek Mishra**

-  Developer Advocate @Yugabyte 
-  Pythonista
-  GDG Chennai, PyCon India
- Socials: @stalwartcoder



```
SELECT * FROM jokes WHERE  
creator='abhishek' AND  
senseofhumor = 'dry' LIMIT 2;
```

- Why did the DBA cross the road?

To get to the other side of the 'view'!

- Why don't we ever see SQL and UML in the same room?

Because they have different views on normalization.

Storms in the Cloud

Does cloud run 24x7?

“Everything fails, all the time”.

– Werner Vogels

Going cloud native? Fasten your seat belts!

- Zone-level outages are commonplace
- Region-level outages become the new normal
- Starting 2011, AWS alone had a major outage *at least once* a year:
 - It can take 4+ hours to recover

AWS Post-Event Summaries

Free AWS Training | Advance your career with AWS Cloud Practitioner Essentials—a free, six-hour,

AWS Post-Event Summaries

The following is a list of post-event summaries from major service events that impacted AWS service availability:

- Summary of the AWS Service Event in the Northern Virginia (US-EAST-1) Region, December, 10th 2021
- Summary of AWS Direct Connect Event in the Tokyo (AP-NORTHEAST-1) Region, September, 2nd 2021
- Summary of the Amazon Kinesis Event in the Northern Virginia (US-EAST-1) Region, November, 25th 2020
- Summary of the Amazon EC2 and Amazon EBS Service Event in the Tokyo (AP-NORTHEAST-1) Region, August 23, 2019
- Summary of the Amazon EC2 DNS Resolution Issues in the Asia Pacific (Seoul) Region (AP-NORTHEAST-2), November 24, 2018.
- Summary of the Amazon S3 Service Disruption in the Northern Virginia (US-EAST-1) Region, February 28, 2017.
- Summary of the AWS Service Event in the Sydney Region, June 8, 2016.
- Summary of the Amazon DynamoDB Service Disruption and Related Impacts in the US-East Region, September 20, 2015.
- Summary of the Amazon EC2, Amazon EBS, and Amazon RDS Service Event in the EU West Region, August 7, 2014.
- Summary of the Amazon SimpleDB Service Disruption, June 13, 2014.
- Summary of the December 17th event in the South America Region (SA-EAST-1), December 20, 2013.
- Summary of the December 24, 2012 Amazon ELB Service Event in the US-East Region, December 24, 2012.

<https://aws.amazon.com/premiumsupport/technology/pes/>

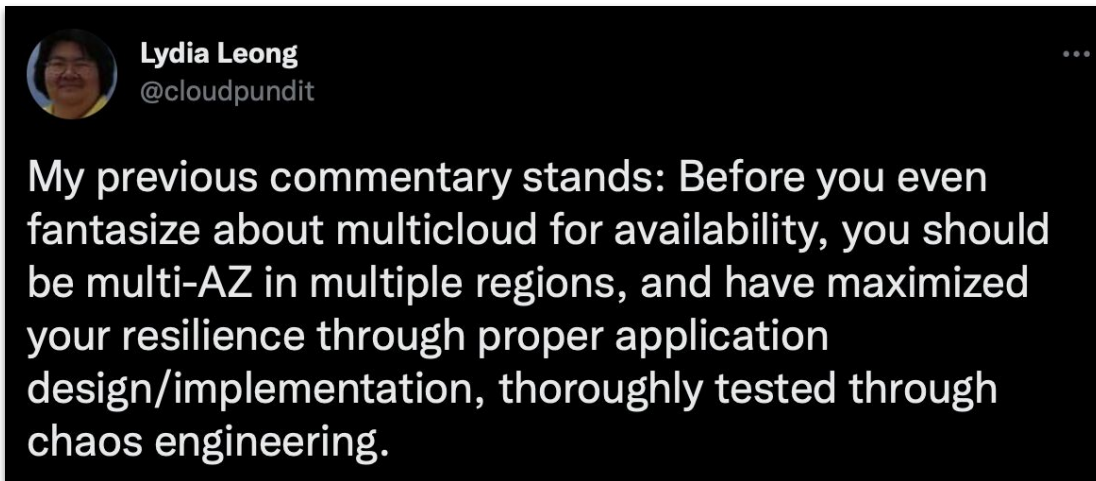
Stuff that is NOT under your control

- Or what causes major cloud outages
- Hardware & networking failures
- Bugs or outages in cloud services:
 - [Amazon Kinesis outage](#) that cascaded to dependent services
- Human errors:
 - [S3 outage](#) caused by a wrong command executed by an AWS team member



Stuff that is under YOUR control

1. Accept the fact that a cloud storm can happen to you
2. Design multi-AZ (at least) and multi-region (better) solutions



<https://twitter.com/cloudpundit>

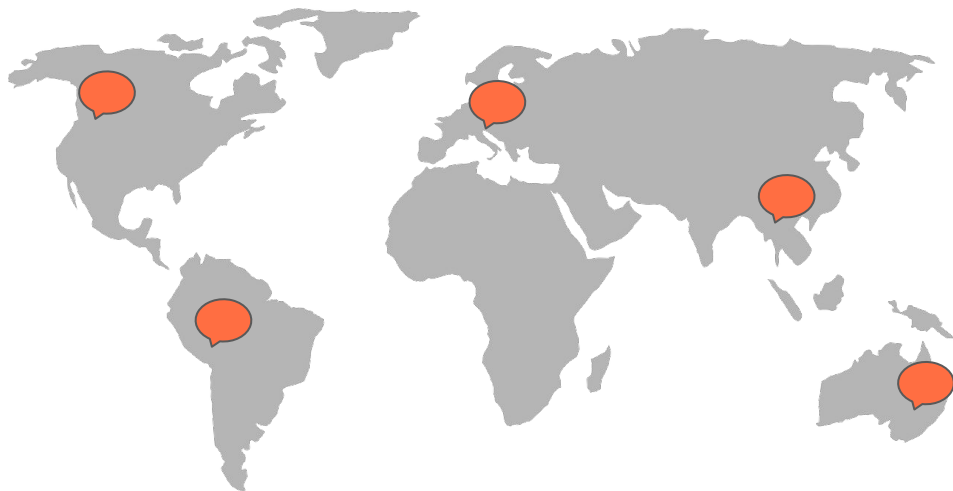
How about speed and latency between zones and regions?

- Zones are interconnected through a high-performance network:
 - GCP: round-trip network latency is under **5 milliseconds** on the 95th percentile
- Regions are interconnected through a regional network ([~latencies for GCP](#)):
 - US West (Oregon)->US East (North Virginia): **60ms**
 - US West (Oregon)->London: **130ms**
 - US West (Oregon)->Mumbai: **220ms**
 - US West (Oregon)->Tokyo: **90ms**

Data Management Patterns for Reliability and Availability

Geo-distributed apps require reliable and HA data layer

Applications that **span multiple geographic locations** for high-availability, compliance and performance needs




Multi-Region Database Deployment With Read Replicas

- Primary database instance is in one geography:
 - Spans multiple zones or local regions
 - Handles both reads and writes
- Read replicas are deployed in other distant locations:
 - Takes over in case of a primary instance failure
 - Serve reads faster for local users



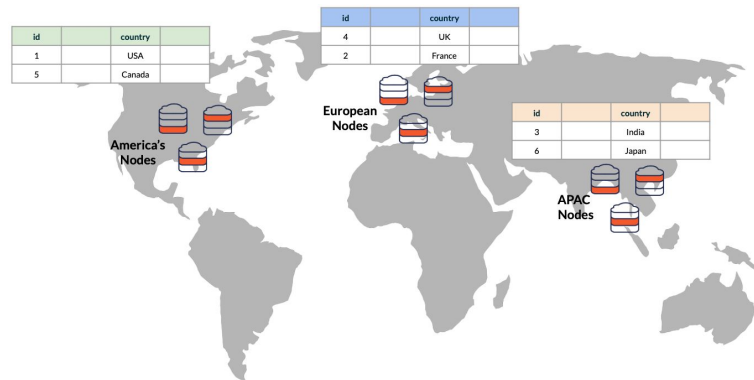
Multiple Standalone Databases With Async Replication

- Each distant geography has its own distributed database instance
- Async replication between standalone databases:
 - Bidirectional and unidirectional
- But synchronously within a single instance



Geo-Partitioned Multi-Region Database Deployment

- Single database instance that spans various geographies
- Data is pinned to a group of nodes from a specific location (NA, EMEA, APAC)
- Low latency for queries within a single geography
 - High latency for cross-region queries



Summary

Learn More?!

- [Architecting for Scale](#) by O'Reilly
- [Designing Data-Intensive Applications](#) by O'Reilly
- Deep dive into a [multi-region DBs deployment options](#)

Exploring Multi-Region Database Deployment Options With a Slack-Messenger

Distributing data, maintaining messaging.



by Denis Magda · Apr. 16, 22 · Database Zone · Tutorial



Liked (3)



Comment (0)





Save



Tweet



3.15K Views



Read more a about incidents:

<https://status.cloud.google.com/summary>

Thank You

Join our community

yugabyte.com/slack

Star the repo



Star

6,752

[yugabyte/yugabyte-db](https://github.com/yugabyte/yugabyte-db)

Try for yourself

cloud.yugabyte.com



yugabyteDB